

Building a Robust Spark Streaming Application



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Understand checkpointing in streaming applications

Understand how driver, executor and receiver fault tolerance works

Build a real world application to work with streaming Twitter data

Build robustness into this application using checkpointing

A Robust Application



Fault tolerance

**The ability of an application to
recover from failures**

A Robust Application



Spark runs on a distributed system

- **Data replication**
- **Process restart on new nodes in case of crashes**




A Robust Application

**Streaming applications
require additional
features to protect
against data loss**


Fault Tolerance with Checkpointing

Checkpointing



Fault tolerance

**Periodically save data to a
reliable storage system**




Checkpointing

Limiting state re-computation in case of failure

Fault tolerance for driver jobs

The Directed Acyclic Graph




A Spark program is split into **discrete tasks**

A task usually has an **RDD** associated with it


The task and its associated RDD **depend** on other tasks

The Directed
Acyclic Graph




The tasks and
RDDs form a
**Directed Acyclic
Graph**

Lineage




An RDD can
always be
reconstructed
using its lineage

Limiting State Re-computations




**Reconstructing an RDD from
scratch is a lot of re-computation**

Limiting State Re-computations



Checkpointing **reduces** how much state has to be recomputed

Limiting State Re-computations



If the driver program crashes it
can recover **from the checkpoint**
rather than the beginning of time

```
sc = SparkContext(appName="StreamingWordCount")
ssc = StreamingContext(sc, 30)

ssc.checkpoint("file:///tmp/spark")
```

Checkpoint Streaming Applications

For production systems, store checkpoints in a reliable replicated source like HDFS or Amazon's S3

Demo

Connect to the Twitter Streaming API to access tweets using the Python tweepy package

Stream tweets to a port on the local machine

Build a Spark application to listen to streaming tweets and perform transformations

Fault Tolerance in Spark Streaming Components

Spark Architecture



Master/Slave architecture


**A central coordinator with many
distributed workers**

Spark Architecture




**You write code, and
submit it to Spark**

Spark Architecture



**A coordinator receives the program
and breaks it into discrete tasks**

Spark Architecture



**Each task is assigned to
worker machines to execute**

Spark Architecture



Driver

**The coordinator process
which executes the user
program**



Executor

**The worker process
responsible for running
individual jobs**



Driver

Converts a user program into physical executables called **tasks**

Schedules tasks on **executors**

Runs in its own Java process

Spark Architecture



Driver



Executor




Executors

Worker processes which run the tasks in a Spark job

Provide **in-memory storage** for **RDDs** so that tasks run close to the data


Each executor is a separate Java process

Spark Architecture



A **cluster manager** for distributed systems launches the driver and executor programs

Spark Architecture



Cluster managers are **pluggable**,
the built-in one is called the
Standalone cluster manager


Spark streaming uses an additional component called **receivers**, one for each input source

Spark Architecture



**Allocate executors to
accept streaming data**

Spark Architecture



**Tasks within those executors
receive data from external sources**

Input Data in Streaming Spark



Receivers

**Collect input data from
different sources**

Receivers



Tasks run **within** executors

Collect data from **input sources** and save them as RDDs

Replicate the collected data to another executor for fault tolerance

Spark Architecture



The streaming application should be able to **recover** from failures of all these components

Driver Fault Tolerance

Driver Fault Tolerance



**Recover from driver crashes
from the last checkpoint state**

Driver Fault Tolerance



No intermediate checkpoint state

- **Set up a brand new state for the job**

Intermediate state present

- **Recreate state from the checkpoint**

Demo

Use the `StreamingContext.getOrCreate()` function to re-create context from stored state

Executor and Receiver Fault Tolerance

Executor Fault Tolerance



**Worker crashes and
recovery are similar to Spark**



Executor Fault Tolerance

All data from external sources are **replicated** on multiple nodes

RDDs from this source data use **lineage** to reconstruct themselves

Receiver Fault Tolerance



Workers running receivers have different fault tolerance mechanics

Receiver Fault Tolerance



Received data is
replicated to
other data nodes

Receiver Fault Tolerance




Sources which are reliable can **resend** data

- HDFS, S3, storage systems

Sources such as Twitter, Kafka may be prone to **some data loss** if the data is **not checkpointed**

Processing Guarantees



Spark streaming provides **exactly
once semantics** for processes



Processing Guarantees

External systems might consume Spark output

Transformed data might get pushed to them **multiple** times

- this has to be accounted for either in Spark or the external system

Overview

Understood the importance of saving intermediate state with checkpointing

Know how driver, worker and receiver fault tolerance works

Built a robust real world application to work with streaming Twitter data